

fren_infrastructure_control_v4 — Scope, Testing & Autoloop Report

Generated 2026-06-11 (autonomous run, Claude Code)

Contents

1	Executive summary	2
2	System architecture	2
3	The agent fleet	2
3.1	Routing graphs	3
4	Runtime services & features	5
5	Retrieval infrastructure	5
6	The improvement loops	6
6.1	1. Fleet judged loop (<code>./run_autoloop.sh</code>)	6
6.2	2. Probe packs (<code>python -m app probe-packs</code>)	6
6.3	3. Proactive context-signal suite (<code>--proactive-probes</code>)	6
6.4	4. Retrieval suite (<code>--retrieval-probes</code>)	6
6.5	5. Branch tests + step contracts	8
6.6	6. Delivery gate policy loop (<code>python -m app improve-gate</code>)	8
6.7	7. RALF end-to-end smoke (<code>python -m app ralf-smoke</code>)	8
6.8	8. Framework primitives behind the loops	8
7	Testing & measured results	12
7.1	Unit suites	12
7.2	Live verifications (this run)	12
7.3	Fleet sweep scoreboard (first full judged baseline)	12
8	Operations record	12
9	Known gaps & next steps	12
10	CLI reference	13

1 Executive summary

fren_infrastructure_control_v4 is a clean-room rebuild of the v3 personal assistant (“Twily”) on top of the **OpenCodeCompilerV2** agent framework. The system runs **147 agents across 15 domains**, entirely on private infrastructure: agents execute on a local Qwen3.5-27B (vLLM, this machine, 4xGPU tensor-parallel), while GLM-5.1 — reached exclusively through opencode — acts as teacher (prompt rewriter) and judge.

The defining property of v4 is that **every capability is autoloop-optimizable**: each agent, branch, and runtime policy carries machine-gradable probes, so the system can be re-tuned automatically against any model. This report documents what is supported, how each piece is tested, every improvement loop and its structure, and the measured results of the first full judged fleet sweep (15.5 hours, 157 units, zero infrastructure failures).

Key measured results:

- **Unit tests**: 1,283 framework + 496 v4 tests passing.
- **Fleet sweep**: 157 units (140 agents + 17 branch tests), 0 failed units, **5 promotions** (best candidate beat baseline above threshold 0.75).
- **Authored test suites**: mean score 0.83 (78% of tests ≥ 0.7).
- **Judged corpus probes**: mean ~ 0.3 ; 46% of agents floor at zero on at least one probe — the honest improvement backlog, enumerated per-test.
- **Retrieval ground truth**: canary facts planted in a seeded 20,101-message corpus are retrieved as the top semantic hit (similarity 0.69 over 20k messages); canary transcript is the top chunk hit.

2 System architecture

```
telegram user --> bot (python-telegram-bot, async)
                    | slash commands --> workflows/* agents
                    | free text -----> persona/orchestrator --> subagents
                    v
                    opencode runtime (sessions, tool permissions)
                    | agents = compiled .md (prompt + scoped bash allowlist)
                    v
local vLLM Qwen3.5-27B-AWQ (4xGPU TP, temp 1.0/topP 0.95/topK 20)
GLM-5.1 via opencode (teacher + judge ONLY - never raw API)
                    |
                    v
                    PostgreSQL (pgvector) - 30+ tables
                    scheduler (41 crons) - checker - web dashboard (HTMX)
```

Two repositories:

- **OpenCodeCompilerV2** (the framework, $\rightarrow$ PyPI `open-agent-compiler`): agent/tool/test models (Pydantic), compiler (agents $\rightarrow$ opencode artifacts), improvement loops, evaluators, judges, workflow DAG executor, in-process interactive runner.
- **fren_infrastructure_control_v4** (the consumer): agent definitions, tools (70+ script tools), telegram/scheduler/web services, probe suites, corpus seeders.

The split-profile mechanism promotes per-model_class (default / fast / analytical / vision), so switching local models re-tunes prompts without touching code — the property the autoloops guarantee.

3 The agent fleet

147 agents in 15 domains. Orchestrators (orange in the graphs) route to specialist subagents; every tool-carrying agent has a per-script bash allowlist (deny-all default, e.g. exactly `python scripts/food_manager.py*`).

domain	agents	what it does (examples)
workflows	37	slash-command entry points: /brief, /goal, /council, RALF loop (5 agents), techtree, shopping, meal_plan, memory
support	26	daily_briefer, document_analyst, stt/tts, email/calendar, image analysis, night_analyst, master_investigator
goals	18	priority_orchestrator (morning/evening), task_triage, nudge_strategist, periodic_checker, winddown, strategy loop
persona	14	orchestrator → quick_ack / responding / thinking; topic_synthesizer, thought_forger, twily_selfie/videographer
research	8	orchestrator → video_fetcher (YT transcripts), website_checker, topic_analyst (prism analysis), price_checker
profile	7	self-examination loop: pattern_observer → hypothesis_generator → validator → discovery_manager → knowledge_compiler
server	7	hardware_agent, filesystem_agent, sessions_agent, vision_analyzer, camera_capture_agent
food	6	orchestrator, recipe_parser, restaurant_intake, meal_planner, product_image_indexer
rp	6	game_master, narrator, character_speaker, adventure_generator, world_editor/updater
vis_simulation	5	scenario_generator → conversation_simulator → quality_scorer loop
funfact	4	data_gatherer → pattern_analyst → fact_writer
workflow_master	4	agent_visualizer, documentation_generator, quality_reviewer
study	3	question_master → answer_grader; session_planner (anti-hallucination source: contracts)
investigation	1	youtube_scout
retrieval	1	fast_retrieval — the universal retriever (see §5)

3.1 Routing graphs

Edges extracted from the live registry (full-id references + same-domain bare-name matches): **79 verified routing edges**. Orchestrator/dispatcher nodes are highlighted.

Cross-domain routing (edge = N references)




Figure 1: Cross-domain routing overview




Figure 2: persona/* — the conversational core

4 Runtime services & features

- **Telegram bot** — free-text chat through **persona/orchestrator** (trivial → quick_ack; substantive → context analysis → responding), photo-caption dispatch, voice (STT/TTS). **11 slash commands**: /help /brief /memory /analyse /invoice /techtree /goal /council /adventure /ralf /study.
- **Delivery contract** — agents never print to the user; they deliver via **emit_guidance.py**. The **delivery gate** (autoloop-tuned policy **policy:delivery_gate**) suppresses duplicates (similarity 0.78, lookback 8), no-op pings, and internal-mechanics leaks at the single send_message chokepoint.
- **Scheduler** — **41 cron entries** (morning/evening planning, 5-minute periodic checks 8–21h, daily profile analysis, weekly conclusion merge, night_analyst, relationship initiator/reflector, topic synthesis, ...).
- **Web dashboard** (HTMX) — Mind tab (mood 5D, vibe, interests), Life tab (goals tree, todos, habits, Eisenhower priorities), run traces with full LLM audit (prompts/thinking/tokens/p50/p95/fallback rate), events charts, images lightbox, artifacts gallery, health.
- **Safety rails** — capture kill-switch (FREN_DISABLE_CAPTURE), X-proof autoloop env, **GPU power guard** (autoloop refuses to start above 200W caps — see §8), prod-DB isolation for all loops.

5 Retrieval infrastructure

One unified entry point (**fetch-context**) fans out to six sources in parallel (3s/source, 15s total), weights, dedupes, ranks:

source	store	weight
context pins	context_pins	1.5
artifact cache	context_cache	1.3
memories	memories (pgvector)	1.2
digital journal	telegram-log service on orp (orange pi, 192.168.0.80:5050)	1.1
embedding chunks	embedding_chunks (transcripts, documents)	1.0

source	store	weight
chat history	<code>chat_messages</code> (pgvector)	1.0

How Twily retrieves: any agent spawns `retrieval/fast_retrieval` (read-only, JSON-only, <15s contract, all six tools). Self-examination uses `session-inspector` (opencode session store: find-by-text/time, session trees) plus the `profile/*` analysis cycle. YouTube transcripts come from SearchAPI.io (`youtube_fetcher.py`), stored full-text + chunk-embedded.

The seeded test corpus (`python -m app seed-retrieval`): 19,799 real v3 messages (Feb–Jun 2026, source DB opened read-only) + 200 full YT transcripts + **10 canary facts** + **1 canary transcript**, embedded with the production encoder. The seeder refuses any non-`autoloop` database. Verified: canary = top-1 hit over 20k messages.

6 The improvement loops

All loops share one engine — the framework’s `IterativeLoop`:

```
baseline --> mutate (teacher rewrites prompt against concrete failures)
      ^           |
      |           v
promote <-- evaluate (run agent on probes → graded scores)
      (if winner >= threshold; snapshot + content-hash lineage)
```

Evaluation always produces `pass_rate`, `score_floor`, `score_mean`, and per-test `score_floor:by_name:<test>` so a failing probe is identifiable by name. Promotion is per-`model_class` (split profiles). Every loop runs against an isolated copy of the production DB (`run_autoloop.sh`), never prod.

6.1 1. Fleet judged loop (`./run_autoloop.sh`)

Every agent is improvable even without authored tests: a **role-fulfilment judge test** is synthesized from the agent’s own definition, merged with its corpus pack and any authored suites. The candidate runs as a REAL opencode session on qwen (tools live, thinking on); the GLM judge grades the outcome (for delivery agents: the *emitted payload*, with hard 0 for non-delivery); the GLM teacher rewrites the prompt against the named failures. 1 round = baseline + mutations for 157 units ~ 15.5h at workers=10.

6.2 2. Probe packs (`python -m app probe-packs`)

Offline generation: per-agent slices of the REAL v3 corpus (domain keyword sampling, EN+PL) → GLM writes self-contained probes with judge rubrics → persisted JSON (137 packs committed). Loaded via the framework’s `ProbePack` primitive. Example probe (persona/orchestrator): “*Ok, please mark the groceries shopping as done, and also remind me to call the dentist tomorrow morning...*” → rubric: must detect three intents.

6.3 3. Proactive context-signal suite (`--proactive-probes`)

For the agents that *initiate* contact (`nudge_strategist`, `periodic_checker`, `winddown`, `event_extractor`): variety / anti-repetition / groundedness / skip probes built from real context snapshots, plus **stale-state replay probes** (resolved-issue scenarios — the agent must NOT re-nudge about a stocked medication, a moved appointment, a resolved call). This suite exists because v3’s engagement collapsed 6x from proactive repetition.

6.4 4. Retrieval suite (`--retrieval-probes`)

Five categories over the seeded corpus (see §5): 10 exact canary probes (`FactRecallEvaluator` — graded recall + zero-on-fabrication), 4 open-ended synthesis probes (groundedness judge), 2 seeded-transcript probes (forces the chunk-search path), 1 live journal probe (orp), 2 self-examination probes. Targets `retrieval/fast_retrieval` (19 tests) and `persona/responding` (6 end-to-end through the delivery contract).




Figure 3: goals/* — planning & proactive agents



Figure 4: RALF loop — plan/execute/verify cycle




Figure 5: research/*

6.5 5. Branch tests + step contracts

Orchestrator dispatch chains are tested deterministically: authored **BranchTests** mock every subagent and assert the **dispatch path order**; **StepContracts** additionally assert per-step inputs (what the orchestrator forwarded) and outputs. 23 branch tests, all domains; 17 ran as fleet units. Example: `study::question-then-grade` asserts `question_master` → `answer_grader` with verbatim `source`: grounding forwarded.

6.6 6. Delivery gate policy loop (python -m app improve-gate)

The runtime send-policy (dedup threshold, lookback, leak/no-op regexes) is a **component**, optimized by `optimize_callable` over 35 replayed real cases — the framework’s “wrap any callable” autoresearch loop with field mutators (`NumericFieldMutator` on `dedup_similarity`, etc.).

6.7 7. RALF end-to-end smoke (python -m app ralf-smoke)

One production-faithful full cycle of the RALF loop (planning → plan evaluation → execution → step evaluation, self-chaining via detached spawns) on the seeded DB. The smoke question requires **two facts from two different stores** (chat canary + transcript canary); the collected outputs are graded with `FactRecallEvaluator`. Exit codes: 0 = completed + recalled; 1 = missed; 2 = timeout.

6.8 8. Framework primitives behind the loops

`IterativeLoop` / `optimize_callable` (autoresearch over any callable), `ComponentRegistry` (content-hash lineage), `OptimisationCriterion` (hard `pass_rate` vs graded `score_floor`), mutators (LLM prompt rewriter, field mutators), `contract_gate`, `ProbeCache`, split profiles, `provider_guard` (AST scan banning raw endpoint literals — `z.ai` is only reachable through `opcode`), workflow DAG executor (slow workflows as autoloop targets), in-process interactive runner (quick agents as library calls — live-verified on qwen: tool loop 1.4s, structured output 1.0s).

profile/* routing




Figure 6: `profile/*` — self-examination

server/* routing




Figure 7: server/*

food/* routing




Figure 8: food/*

7 Testing & measured results

7.1 Unit suites

- Framework: **1,283 tests passing** (models, compiler, evaluators incl. FactRecall + ProbePack, loops, DAG executor, runner, step contracts).
- v4: **496 passing, 4 skipped, 1 xfailed** (tools mocked per design, domains, branch contracts, runtime, gate policy, parity).

7.2 Live verifications (this run)

- Canary retrieval: top-1 semantic hit over 20,101 messages (0.69); canary transcript: top-1 chunk hit.
- Interactive runner on qwen: tool-call loop 1.4s; structured output 1.0s (single leading system message — strict chat templates reject a second).
- Blocked-call forensics over 37k session messages → two scoped fixes (read-only `ls` orientation tool; `uv run` alias for an agent’s own script).

7.3 Fleet sweep scoreboard (first full judged baseline)

157 units, 0 failures, 15.5h, GPUs capped 200W, zero platform incidents.

metric	value
promoted	5 (topic_synthesizer 0.92, stt_processor 0.85, documentation_generator 0.78, branch:techtree_orchestrator 0.75, branch:study/question_master)
authored tests	mean 0.83, 78% ≥ 0.7 , 10% zeros
role-fulfilment judge	mean 0.31, 27% ≥ 0.7 , 46% zeros
corpus pack probes (650)	mean 0.26, 19% ≥ 0.7 , 46% zeros

Reading: the summary “winner score” is the **floor** across an agent’s tests, so one zeroed probe floors the unit. The per-test metrics (`score_floor:by_name:...` in every snapshot) enumerate exactly which probe fails for which agent — the round-2 worklist. Notable anomaly: 2 agents (persona/responding, profile/knowledge_compiler) ace the role test but zero all pack probes (delivery-contract interaction under corpus-style prompts) — under active improvement in the retrieval loop at report time.

8 Operations record

- **Host stability root cause** (confirmed by controlled A/B): synchronized multi-GPU power transients through a degrading 12V path. 8 concurrent inference streams: hard-lock in 23s at 250W caps; clean at 150W and 200W (5-minute tests); then 15.5h sustained at 200W with only sporadic corrected link errors and zero USB resets. `run_autoloop.sh` now **refuses to start above 200W caps** (override: `FREN_AUTOLOOP_MAX_PL` after hardware repair). Hardware to-do: reseal/replace 12V connectors / PSU coupling; silicon fine.
- **Sampling**: qwen at temperature 1.0 / topP 0.95 / topK 20 (opencode.json); GLM teacher 0.4–0.6, judge 0.0–0.3 (determinism by design); vLLM `--generation-config auto` as the fallback floor.
- **Isolation**: every loop runs on `<db>_autoloop` (full prod copy, plus the seeded retrieval corpus), with capture kill-switch and X-stripped env by default (`--with-capture` is an explicit opt-in).

9 Known gaps & next steps

1. **Round-2 deeper pass** on the 46%-floor agents (worklist = per-test metrics from this sweep); rounds=2–3 on the weakest ~20 first.

2. **Retrieval + RALF smoke results** — in flight at report time; will land in snapshots/.oac and the next report revision.
3. **Model-switch drill** (task #206): retune against glm-4.7-as-target and diff per-model_class promotions — the final proof of model portability.
4. **Proactive context parity** (task #198): feed v4 proactive agents v3's exact context shape.
5. **Bot watchdog** (task #200): durable in-process hang detector.
6. Cleanup: dormant `run_agent_direct` raw-z.ai path (unused in prod) to be deleted or repointed at vLLM.

10 CLI reference

```
./run_autoloop.sh [--with-capture] [--refresh] <improve args>
python -m app improve --list | --proactive-probes | --retrieval-probes \
    [--agent id]... [--rounds N] [--workers N]
python -m app probe-packs [--agents ...] [--domains ...] [--refresh]
python -m app seed-retrieval [--no-embed] [--limit N] [--yt-limit N]
python -m app ralf-smoke [--timeout-min 25] [--question "..."]
python -m app improve-gate [--rounds 4] [--no-promote]
python -m app bot | scheduler | checker | web | compile
```